

Incident Severity Guide

wiki/sre/incident-severity · Updated 2026-03-28 · Owner: SRE

Purpose

This guide defines how Northwind classifies incident severity so that response matches impact. It is maintained by the SRE team and is referenced by the on-call runbook. Classify every incident at the start, and adjust the classification as you learn more.

A shared severity language lets everyone understand how serious an incident is without a long explanation, which speeds up the response.

How to classify

Classify by customer impact, not by internal cause. A failing internal component that customers do not notice is lower severity than a small but customer-visible failure. When in doubt, start higher and downgrade.

Severity levels

- SEV1: full outage or data loss affecting most customers; page immediately.
- SEV2: major feature broken or significant degradation.
- SEV3: minor issue with a workaround.

SEV1

A SEV1 is an all-hands situation. Open an incident channel immediately, assign an incident commander, and page the service owner. Communicate frequently with clear status updates. The priority is restoring service, not finding the root cause. Data loss is always at least a SEV1.

SEV2

A SEV2 is serious but not a full outage. A major feature may be broken or performance badly degraded for many users. Respond promptly during the day and page out of hours if customer impact is ongoing. Many SEV2s become SEV1s if left unaddressed, so do not sit on them.

SEV3

A SEV3 is a minor issue with a reasonable workaround. Track it and fix it during normal working hours; it does not require paging. If a SEV3 starts affecting more customers, raise it.

Classifying and adjusting

Classify at the start of an incident and adjust as you learn more. It is better to start high and downgrade than to under-call a serious incident. If customer impact grows, raise the severity and expand the response.

Communication expectations

- SEV1: status updates every 30 minutes to stakeholders.
- SEV2: regular updates until the issue is resolved.
- SEV3: a tracked ticket with normal updates.

Who communicates

The incident commander owns communication for SEV1 and SEV2 incidents, or delegates it to a dedicated communications lead for large incidents. Responders should not be interrupted for status; the commander aggregates and shares it.

Postmortems

SEV1 and SEV2 require a written postmortem within five business days. Postmortems are blameless and focus on contributing factors and actions that reduce the chance of recurrence. Share them widely so the whole organisation learns.

Blameless culture

A blameless postmortem assumes everyone acted reasonably given what they knew. The goal is to fix the system, not to assign fault. People report problems honestly only when they are not punished for them.

Worked example

A payment feature starts failing for a subset of customers. The on-call engineer classifies it as a SEV2, opens an incident, and assigns a commander. As more customers are affected, they raise it to SEV1 and page the service owner. After mitigation, they hold a blameless postmortem within five business days and track the follow-up actions to completion.

Frequently asked questions

Is data loss always a SEV1? Yes, at minimum.

Should I classify by cause or impact? By customer impact.

What if I am not sure of the severity? Start higher and downgrade as you learn more.

Common mistakes

- Under-calling a serious incident to avoid waking people up.
- Classifying by internal cause instead of customer impact.
- Treating a postmortem as a search for someone to blame.
- Leaving a SEV2 unattended until it becomes a SEV1.

Roles

The incident commander coordinates the response and owns communication. Responders focus on diagnosis and mitigation. A scribe captures the timeline for the postmortem.

Related documents

See the On-Call Runbook for how incidents are handled and the Monitoring and Alerting Setup for the alerts that surface them.

Examples by severity

Concrete examples make the levels easier to apply under pressure. Use these as a guide, and when a situation sits between two levels, choose the higher one.

- SEV1: checkout is down for all customers; the primary database is unreachable; customer data has

been exposed or lost.

- SEV2: payments fail for a region; a core page is badly degraded; a key integration is down with no

workaround.

- SEV3: a non-critical report is delayed; a cosmetic defect; an issue affecting a handful of users

with a workaround.

Severity and response times

Severity drives how quickly and how widely we respond. These expectations exist so that effort matches impact and nobody has to negotiate the response in the middle of an incident.

- SEV1: immediate response, day or night, with an incident commander and updates every 30 minutes.
- SEV2: prompt response during the day, paged out of hours if customer impact continues.

- SEV3: handled during normal working hours as tracked work.

Downgrade only when customer impact has genuinely reduced, and record the change so the timeline is accurate.