

On-Call Runbook

wiki/sre/oncall-runbook · Updated 2026-09-12 · Owner: SRE

Overview

This runbook describes how on-call works for Northwind services and what to do when you are paged. It is maintained by the SRE team. Read it before your first rotation and keep it within reach while you are on call.

On-call exists to protect customers. The person on call is empowered to take whatever reasonable action is needed to restore service, including rolling back releases and escalating to anyone in the company.

Before your rotation

Prepare before your week starts so you are not scrambling when the first page arrives.

- Confirm you can log in to the dashboards and the alerting console.
- Confirm you can silence and acknowledge alerts.
- Install and test the paging app on your phone.
- Read the recent incident history and any known risks.

The rotation

The on-call engineer owns the pager for one week. Hand-off happens at a fixed time with a short sync covering open incidents, risky changes, and

anything to watch.

- Acknowledge pages within 5 minutes.
- Escalate to the secondary on-call if a page is unacknowledged after 15 minutes.
- Keep your phone charged and your laptop reachable during your rotation.

Hand-off

At the start and end of each rotation, run a short hand-off. The outgoing engineer summarises open incidents, recent deploys, and anything fragile. The incoming engineer confirms access and asks questions.

Responding to a page

When you are paged, first decide whether customers are affected.

- Confirm the alert is real and not a monitoring glitch.
- Classify the severity using the incident severity guide.
- For a production outage, open an incident channel and assign an incident commander.
- Post status updates every 30 minutes until the incident is resolved.

During an incident

Keep communication clear and frequent. The incident commander coordinates; everyone else focuses on their task. Prefer fast, reversible

actions such as rolling back a recent release over risky forward fixes. Capture a rough timeline as you go so the postmortem is easier later.

Mitigate first, diagnose later

The first priority is restoring service, not understanding why it broke. If a recent deploy is the likely cause, roll it back before investigating. Diagnosis can happen once customers are no longer affected.

Common actions

- Roll back the most recent release if the incident started right after a deploy.
- Pin Cloud Run traffic to the last known-good revision.
- Scale up if the service is saturated.
- Fail over to a healthy region if one region is degraded.

Communication

For a customer-facing outage, keep stakeholders informed with regular updates even if there is no news. Silence makes people assume the worst. A short "still investigating, next update in 30 minutes" is enough.

After the incident

Once service is restored, close the incident channel and schedule a postmortem for SEV1 and SEV2 incidents. Postmortems are blameless and focus on systems and process, not individuals.

The postmortem

A good postmortem captures the timeline, the contributing factors, what went well, and concrete actions to reduce the chance of recurrence. Assign owners and due dates to the actions, and track them to completion.

Worked example

At 2 a.m. the on-call engineer is paged for a rising error rate. They acknowledge within two minutes, confirm the alert is real, and classify it as a SEV2. The dashboards show the errors began right after a deploy, so they pin Cloud Run traffic to the previous revision. Errors recover. They post an update, open an incident, and schedule a postmortem for the next morning.

Frequently asked questions

What if I cannot fix it? Escalate. Pull in the service owner, the secondary on-call, or your manager. Escalating early is good judgement.

Do I diagnose or mitigate first? Mitigate first. Restore service, then investigate.

How often do I update stakeholders? Every 30 minutes for a SEV1, and regularly for a SEV2, even if there is no change.

Common mistakes

- Trying to debug a bad release in production instead of rolling back.
- Going quiet during an incident so stakeholders assume the worst.

- Not writing down the timeline, making the postmortem guesswork.
- Hesitating to escalate because it feels like admitting failure.

Escalation

If you are unsure or overwhelmed, escalate early. Pull in the service owner, the secondary on-call, or your manager. Escalating is never a failure; a prolonged customer-impacting outage is.

Related documents

See the Incident Severity Guide for classification, the Monitoring and Alerting Setup for dashboards and alerts, and the CI/CD Pipeline Standards for how to roll back a release.

Handover notes

A good handover is the difference between a smooth week and a chaotic one. At the end of your rotation, write a short note for the next engineer covering anything they should watch.

- Open incidents and their current status.
- Recent or upcoming deploys that carry risk.
- Any alert that has been firing and why.
- Temporary silences you set and when they expire.

Keep the note in the agreed place so it is easy to find, and walk through it live at hand-off so the incoming engineer can ask questions.